

Computational Bootcamp 2: Intro to R

Ankushi Mitra

Department of Government
Georgetown University

August 15, 2023

What We'll Be Covering Overall

- ① Software installation, file management
- ② Basics of R: data structures, writing code, creating objects, packages
- ③ R: working with datasets
- ④ More R: data cleaning, visualization
- ⑤ LaTeX: producing documents with Markdown and Overleaf

What We'll Be Covering Today

① Basic Steps of Data Analysis

What We'll Be Covering Today

- ① Basic Steps of Data Analysis
- ② The R Studio Interface

What We'll Be Covering Today

- ① Basic Steps of Data Analysis
- ② The R Studio Interface
- ③ Data Structures in R

What We'll Be Covering Today

- ① Basic Steps of Data Analysis
- ② The R Studio Interface
- ③ Data Structures in R
- ④ Working with Objects in R

What We'll Be Covering Today

- ① Basic Steps of Data Analysis
- ② The R Studio Interface
- ③ Data Structures in R
- ④ Working with Objects in R
- ⑤ Packages in R

Basic Steps of Data Analysis

- Specify research question

Basic Steps of Data Analysis

- Specify research question
- Develop research design

Basic Steps of Data Analysis

- Specify research question
- Develop research design
- Collect data

Basic Steps of Data Analysis

- Specify research question
- Develop research design
- Collect data
- Preprocess data

Basic Steps of Data Analysis

- Specify research question
- Develop research design
- Collect data
- Preprocess data
 - The data you have can contain errors, be incomplete, may need additional variable construction, etc.

Basic Steps of Data Analysis

- Specify research question
- Develop research design
- Collect data
- Preprocess data
 - The data you have can contain errors, be incomplete, may need additional variable construction, etc.
 - Clean the dataset so that it can be used for analysis

Basic Steps of Data Analysis

- Specify research question
- Develop research design
- Collect data
- Preprocess data
 - The data you have can contain errors, be incomplete, may need additional variable construction, etc.
 - Clean the dataset so that it can be used for analysis
 - Summarizing/Visualizing data

Basic Steps of Data Analysis

- Specify research question
- Develop research design
- Collect data
- Preprocess data
 - The data you have can contain errors, be incomplete, may need additional variable construction, etc.
 - Clean the dataset so that it can be used for analysis
 - Summarizing/Visualizing data
 - Statistical analysis

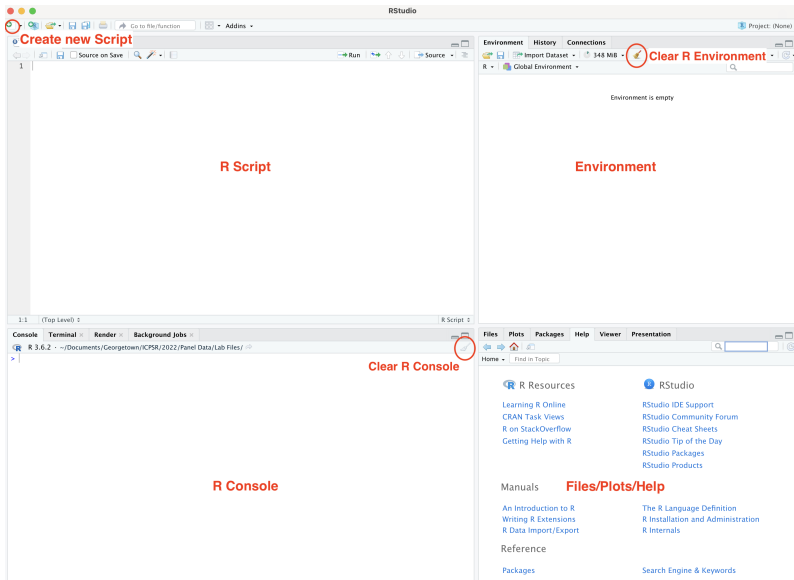
Basic Steps of Data Analysis

- Specify research question
- Develop research design
- Collect data
- Preprocess data
 - The data you have can contain errors, be incomplete, may need additional variable construction, etc.
 - Clean the dataset so that it can be used for analysis
 - Summarizing/Visualizing data
 - Statistical analysis
 - Summarizing/Visualizing results

Basic Steps of Data Analysis

- Specify research question
- Develop research design
- Collect data
- Preprocess data
 - The data you have can contain errors, be incomplete, may need additional variable construction, etc.
 - Clean the dataset so that it can be used for analysis
 - Summarizing/Visualizing data
 - Statistical analysis
 - Summarizing/Visualizing results

R Studio Interface



R Studio Interface

- R Script
 - Write your code here

R Studio Interface

- R Script
 - Write your code here
 - Each line is a separate command

R Studio Interface

- R Script
 - Write your code here
 - Each line is a separate command
 - "Run" a line by highlighting it and clicking "Run" or pressing Cmd+Return (Mac) or Ctrl+Return (Windows)

R Studio Interface

- R Script
 - Write your code here
 - Each line is a separate command
 - "Run" a line by highlighting it and clicking "Run" or pressing Cmd+Return (Mac) or Ctrl+Return (Windows)
 - Code should be written so humans can easily understand what's happening. Annotate/Comment your code by prefacing a line with `#`

R Studio Interface

- R Script
 - Write your code here
 - Each line is a separate command
 - "Run" a line by highlighting it and clicking "Run" or pressing Cmd+Return (Mac) or Ctrl+Return (Windows)
 - Code should be written so humans can easily understand what's happening. Annotate/Comment your code by prefacing a line with #
 - The script can be saved as a .R file

R Studio Interface

- R Script
 - Write your code here
 - Each line is a separate command
 - "Run" a line by highlighting it and clicking "Run" or pressing Cmd+Return (Mac) or Ctrl+Return (Windows)
 - Code should be written so humans can easily understand what's happening. Annotate/Comment your code by prefacing a line with #
 - The script can be saved as a .R file
- Console
 - Where results are displayed

R Studio Interface

- R Script
 - Write your code here
 - Each line is a separate command
 - "Run" a line by highlighting it and clicking "Run" or pressing Cmd+Return (Mac) or Ctrl+Return (Windows)
 - Code should be written so humans can easily understand what's happening. Annotate/Comment your code by prefacing a line with `#`
 - The script can be saved as a .R file
- Console
 - Where results are displayed
 - You can type code directly into the console **but this is bad practice since the Console input cannot be saved**

R Studio Interface

- R Script
 - Write your code here
 - Each line is a separate command
 - "Run" a line by highlighting it and clicking "Run" or pressing Cmd+Return (Mac) or Ctrl+Return (Windows)
 - Code should be written so humans can easily understand what's happening. Annotate/Comment your code by prefacing a line with `#`
 - The script can be saved as a .R file
- Console
 - Where results are displayed
 - You can type code directly into the console **but this is bad practice since the Console input cannot be saved**
- Environment
 - Where objects are stored

R Studio Interface

- R Script
 - Write your code here
 - Each line is a separate command
 - "Run" a line by highlighting it and clicking "Run" or pressing Cmd+Return (Mac) or Ctrl+Return (Windows)
 - Code should be written so humans can easily understand what's happening. Annotate/Comment your code by prefacing a line with `#`
 - The script can be saved as a .R file
- Console
 - Where results are displayed
 - You can type code directly into the console **but this is bad practice since the Console input cannot be saved**
- Environment
 - Where objects are stored

Object Class/Type/Mode

- character: non-numeric (text)

Object Class/Type/Mode

- character: non-numeric (text)
- numeric: numbers

Object Class/Type/Mode

- character: non-numeric (text)
- numeric: numbers
 - Includes "integer" and "double" — concerns level of precision; R works this out behind the scenes

Object Class/Type/Mode

- character: non-numeric (text)
- numeric: numbers
 - Includes "integer" and "double" — concerns level of precision; R works this out behind the scenes
- logical: TRUE and FALSE

Object Class/Type/Mode

- character: non-numeric (text)
- numeric: numbers
 - Includes "integer" and "double" — concerns level of precision; R works this out behind the scenes
- logical: TRUE and FALSE
- factor: a numeric which R has categorized into "levels"

Object Class/Type/Mode

- character: non-numeric (text)
- numeric: numbers
 - Includes "integer" and "double" — concerns level of precision; R works this out behind the scenes
- logical: TRUE and FALSE
- factor: a numeric which R has categorized into "levels"

Data Structures

- Vectors
 - Data analysis is not possible without data structures for data. R has several important data structures that shape how information is stored and processed.

Data Structures

- Vectors
 - Data analysis is not possible without data structures for data. R has several important data structures that shape how information is stored and processed.
 - *Vectors* hold a sequence of values of the same data type, like numbers, characters, or logical values. They are one-dimensional arrays or lists of values that contain one and only one type of data. Vector types in R include: logical, integer, double, character, complex, raw.

Data Structures

- Vectors
 - Data analysis is not possible without data structures for data. R has several important data structures that shape how information is stored and processed.
 - *Vectors* hold a sequence of values of the same data type, like numbers, characters, or logical values. They are one-dimensional arrays or lists of values that contain one and only one type of data. Vector types in R include: logical, integer, double, character, complex, raw.
- Matrices
 - *Matrices* are two-dimensional data structures in R that consists of rows and columns of values. Every element is of the same type. It's essentially a collection of vectors of the same length, organized in rows and columns.

Data Structures

- Vectors
 - Data analysis is not possible without data structures for data. R has several important data structures that shape how information is stored and processed.
 - *Vectors* hold a sequence of values of the same data type, like numbers, characters, or logical values. They are one-dimensional arrays or lists of values that contain one and only one type of data. Vector types in R include: logical, integer, double, character, complex, raw.
- Matrices
 - *Matrices* are two-dimensional data structures in R that consists of rows and columns of values. Every element is of the same type. It's essentially a collection of vectors of the same length, organized in rows and columns.
 - Most data contains at least numeric information and character information. That is why we don't usually use matrices.

Data Structures

- Data Frames
 - *Data Frames* and *Tibbles* are what you will usually use. They are a tabular data structure in R, similar to a spreadsheet or a database table. They can store different types of data (e.g., numeric, character) in columns. Each column is a vector, and all columns must have the same length.

Data Structures

- Data Frames
 - *Data Frames* and *Tibbles* are what you will usually use. They are a tabular data structure in R, similar to a spreadsheet or a database table. They can store different types of data (e.g., numeric, character) in columns. Each column is a vector, and all columns must have the same length.
 - Most times, each column will be of one type while a given row will contain many different types. We usually refer to columns as *variables* and rows as *observations*.

Data Structures

- Data Frames
 - *Data Frames* and *Tibbles* are what you will usually use. They are a tabular data structure in R, similar to a spreadsheet or a database table. They can store different types of data (e.g., numeric, character) in columns. Each column is a vector, and all columns must have the same length.
 - Most times, each column will be of one type while a given row will contain many different types. We usually refer to columns as *variables* and rows as *observations*.
- Example
 - Inspect R's built-in dataset, *mtcars*

Objects and Assignment

- In R, to work with values/data, we *assign* them to *objects*

Objects and Assignment

- In R, to work with values/data, we *assign* them to *objects*
 - Object: named "box" or "container" to store values/data. Think of it as a box where you can store numbers, text, or other types of information. You can give the box a name, and then you can use that name to access and manipulate the data inside.

Objects and Assignment

- In R, to work with values/data, we *assign* them to *objects*
 - Object: named "box" or "container" to store values/data. Think of it as a box where you can store numbers, text, or other types of information. You can give the box a name, and then you can use that name to access and manipulate the data inside.
 - Using the assignment operator: `<-`

Objects and Assignment

- In R, to work with values/data, we *assign* them to *objects*
 - Object: named "box" or "container" to store values/data. Think of it as a box where you can store numbers, text, or other types of information. You can give the box a name, and then you can use that name to access and manipulate the data inside.
 - Using the assignment operator: <-
 - "Less than" (Shift+comma) Hyphen

Objects and Assignment

- In R, to work with values/data, we *assign* them to *objects*
 - Object: named "box" or "container" to store values/data. Think of it as a box where you can store numbers, text, or other types of information. You can give the box a name, and then you can use that name to access and manipulate the data inside.
 - Using the assignment operator: <-
 - "Less than" (Shift+comma) Hyphen
 - object <- values/data

Objects and Assignment

- In R, to work with values/data, we *assign* them to *objects*
 - Object: named "box" or "container" to store values/data. Think of it as a box where you can store numbers, text, or other types of information. You can give the box a name, and then you can use that name to access and manipulate the data inside.
 - Using the assignment operator: `<-`
 - "Less than" (Shift+comma) Hyphen
 - `object <- values/data`
- Exercise

Objects and Assignment

- In R, to work with values/data, we *assign* them to *objects*
 - Object: named "box" or "container" to store values/data. Think of it as a box where you can store numbers, text, or other types of information. You can give the box a name, and then you can use that name to access and manipulate the data inside.
 - Using the assignment operator: `<-`
 - "Less than" (Shift+comma) Hyphen
 - `object <- values/data`
- Exercise
 - Write two arithmetic operations in R and assign them to unique names. Then perform arithmetic operations using the named results.

Objects and Assignment

- In R, to work with values/data, we *assign* them to *objects*
 - Object: named "box" or "container" to store values/data. Think of it as a box where you can store numbers, text, or other types of information. You can give the box a name, and then you can use that name to access and manipulate the data inside.
 - Using the assignment operator: `<-`
 - "Less than" (Shift+comma) Hyphen
 - `object <- values/data`
- Exercise
 - Write two arithmetic operations in R and assign them to unique names. Then perform arithmetic operations using the named results.

```
a <- 5 + 5 + 5
```

```
b <- 6 - 6 - 6
```

```
a + b
```


Objects, Functions, and Arguments

- In R, we use *commands* (or *functions*) to perform tasks on these *objects*

Objects, Functions, and Arguments

- In R, we use *commands* (or *functions*) to perform tasks on these *objects*
 - Data analysis requires a lot more than just basic arithmetic. R contains many *functions* that can perform mathematical operations, process data, run analyses, and more.

Objects, Functions, and Arguments

- In R, we use *commands* (or *functions*) to perform tasks on these *objects*
 - Data analysis requires a lot more than just basic arithmetic. R contains many *functions* that can perform mathematical operations, process data, run analyses, and more.
 - A function is like a recipe that tells R to perform a specific task. You give the function some input (called arguments), and it gives you an output based on the task it's designed to do.

Objects, Functions, and Arguments

- In R, we use *commands* (or *functions*) to perform tasks on these *objects*
 - Data analysis requires a lot more than just basic arithmetic. R contains many *functions* that can perform mathematical operations, process data, run analyses, and more.
 - A function is like a recipe that tells R to perform a specific task. You give the function some input (called arguments), and it gives you an output based on the task it's designed to do.
 - Type your function, then your object and other arguments in parentheses. For example: `command(object)`

Objects, Functions, and Arguments

- In R, we use *commands* (or *functions*) to perform tasks on these *objects*
 - Data analysis requires a lot more than just basic arithmetic. R contains many *functions* that can perform mathematical operations, process data, run analyses, and more.
 - A function is like a recipe that tells R to perform a specific task. You give the function some input (called arguments), and it gives you an output based on the task it's designed to do.
 - Type your function, then your object and other arguments in parentheses. For example: `command(object)`
 - Exercise: calculate the average of the 2 objects you created earlier using the `mean()` function.

Objects, Functions, and Arguments

- In R, we use *commands* (or *functions*) to perform tasks on these *objects*
 - Data analysis requires a lot more than just basic arithmetic. R contains many *functions* that can perform mathematical operations, process data, run analyses, and more.
 - A function is like a recipe that tells R to perform a specific task. You give the function some input (called arguments), and it gives you an output based on the task it's designed to do.
 - Type your function, then your object and other arguments in parentheses. For example: `command(object)`
 - Exercise: calculate the average of the 2 objects you created earlier using the `mean()` function.

`mean(a , b)`

Objects, Functions, and Arguments

- R functions typically contain many *arguments*. Your object is usually one argument, other arguments can be thought of as "options". For example, `mean()` has `x`, `trim`, and `na.rm`.

Objects, Functions, and Arguments

- R functions typically contain many *arguments*. Your object is usually one argument, other arguments can be thought of as "options". For example, `mean()` has `x`, `trim`, and `na.rm`.
- Many arguments have default values and don't need to be included in function calls.

Objects, Functions, and Arguments

- R functions typically contain many *arguments*. Your object is usually one argument, other arguments can be thought of as "options". For example, `mean()` has `x`, `trim`, and `na.rm`.
- Many arguments have default values and don't need to be included in function calls.
- Functions are recognizable because they end in `()`. For example: `command(object, other arguments)`

Objects, Functions, and Arguments

- R functions typically contain many *arguments*. Your object is usually one argument, other arguments can be thought of as "options". For example, `mean()` has `x`, `trim`, and `na.rm`.
- Many arguments have default values and don't need to be included in function calls.
- Functions are recognizable because they end in `()`. For example: `command(object, other arguments)`
- Documentation for functions can be easily accessed by prefacing the function name with `?` and dropping the `()`. The documentation typically includes a description, a list of the arguments, references, a list of related functions, and examples. The examples are incredibly useful. For example: `?mean`

Packages

- Opening RStudio automatically loads base R, a fundamental collection of code and functions.

Packages

- Opening RStudio automatically loads base R, a fundamental collection of code and functions.
- R is an extensible programming language. Users contribute millions of lines of code that can be used by other users without condition or compensation. The main mode of contribution are R packages.

Packages

- Opening RStudio automatically loads base R, a fundamental collection of code and functions.
- R is an extensible programming language. Users contribute millions of lines of code that can be used by other users without condition or compensation. The main mode of contribution are R packages.
- A package is a collection of functions, data, and documentations which is publicly shared to enhance the functionality of R.

Packages

- Opening RStudio automatically loads base R, a fundamental collection of code and functions.
- R is an extensible programming language. Users contribute millions of lines of code that can be used by other users without condition or compensation. The main mode of contribution are R packages.
- A package is a collection of functions, data, and documentations which is publicly shared to enhance the functionality of R.

Packages

- Installing packages

Packages

- Installing packages
 - Packages can be installed using `install.packages("")`. Include the name of the desired package in quotes as the only argument to the function.

Packages

- Installing packages
 - Packages can be installed using `install.packages("")`. Include the name of the desired package in quotes as the only argument to the function.
 - Example: Install the package, *tidyverse*

Packages

- Installing packages
 - Packages can be installed using `install.packages("")`. Include the name of the desired package in quotes as the only argument to the function.
 - Example: Install the package, *tidyverse*
 - It is customary to never include `install.packages()` in a .R script.

Packages

- Installing packages
 - Packages can be installed using `install.packages("")`. Include the name of the desired package in quotes as the only argument to the function.
 - Example: Install the package, *tidyverse*
 - It is customary to never include `install.packages()` in a .R script.
- Using packages
 - After installation, packages need to be loaded once per R session using the `library()` function. Include the name of the desired package without quotes as the only argument to the function.

Packages

- Installing packages
 - Packages can be installed using `install.packages("")`. Include the name of the desired package in quotes as the only argument to the function.
 - Example: Install the package, *tidyverse*
 - It is customary to never include `install.packages()` in a .R script.
- Using packages
 - After installation, packages need to be loaded once per R session using the `library()` function. Include the name of the desired package without quotes as the only argument to the function.
 - Example: Load the package, *tidyverse*

Packages

- Installing packages
 - Packages can be installed using `install.packages("")`. Include the name of the desired package in quotes as the only argument to the function.
 - Example: Install the package, *tidyverse*
 - It is customary to never include `install.packages()` in a .R script.
- Using packages
 - After installation, packages need to be loaded once per R session using the `library()` function. Include the name of the desired package without quotes as the only argument to the function.
 - Example: Load the package, *tidyverse*
 - It is a good idea to include `library()` statements at the top of scripts for each package used in the script. This way it is obvious at the top of the script which packages are necessary.