

# Computational Bootcamp 2: Intro to R

Ankushi Mitra

Department of Government  
Georgetown University

August 15, 2023

# What We'll Be Covering Overall

- ① Software installation, file management
- ② Basics of R: data structures, writing code, creating objects, packages
- ③ R: working with datasets
- ④ More R: data cleaning, visualization
- ⑤ LaTeX: producing documents with Markdown and Overleaf

# What We'll Be Covering Today

## ① Basic Steps of Data Analysis

# What We'll Be Covering Today

- ① Basic Steps of Data Analysis
- ② The R Studio Interface

# What We'll Be Covering Today

- ① Basic Steps of Data Analysis
- ② The R Studio Interface
- ③ Data Structures in R

# What We'll Be Covering Today

- ① Basic Steps of Data Analysis
- ② The R Studio Interface
- ③ Data Structures in R
- ④ Working with Objects in R

# What We'll Be Covering Today

- ① Basic Steps of Data Analysis
- ② The R Studio Interface
- ③ Data Structures in R
- ④ Working with Objects in R
- ⑤ Packages in R

# Basic Steps of Data Analysis

- Specify research question



# Basic Steps of Data Analysis

- Specify research question
- Develop research design

# Basic Steps of Data Analysis

- Specify research question
- Develop research design
- Collect data

# Basic Steps of Data Analysis

- Specify research question
- Develop research design
- Collect data
- Preprocess data

# Basic Steps of Data Analysis

- Specify research question
- Develop research design
- Collect data
- Preprocess data
  - The data you have can contain errors, be incomplete, may need additional variable construction, etc.

# Basic Steps of Data Analysis

- Specify research question
- Develop research design
- Collect data
- Preprocess data
  - The data you have can contain errors, be incomplete, may need additional variable construction, etc.
  - Clean the dataset so that it can be used for analysis

# Basic Steps of Data Analysis

- Specify research question
- Develop research design
- Collect data
- Preprocess data
  - The data you have can contain errors, be incomplete, may need additional variable construction, etc.
  - Clean the dataset so that it can be used for analysis
  - Summarizing/Visualizing data

# Basic Steps of Data Analysis

- Specify research question
- Develop research design
- Collect data
- Preprocess data
  - The data you have can contain errors, be incomplete, may need additional variable construction, etc.
  - Clean the dataset so that it can be used for analysis
  - Summarizing/Visualizing data
  - Statistical analysis

# Basic Steps of Data Analysis

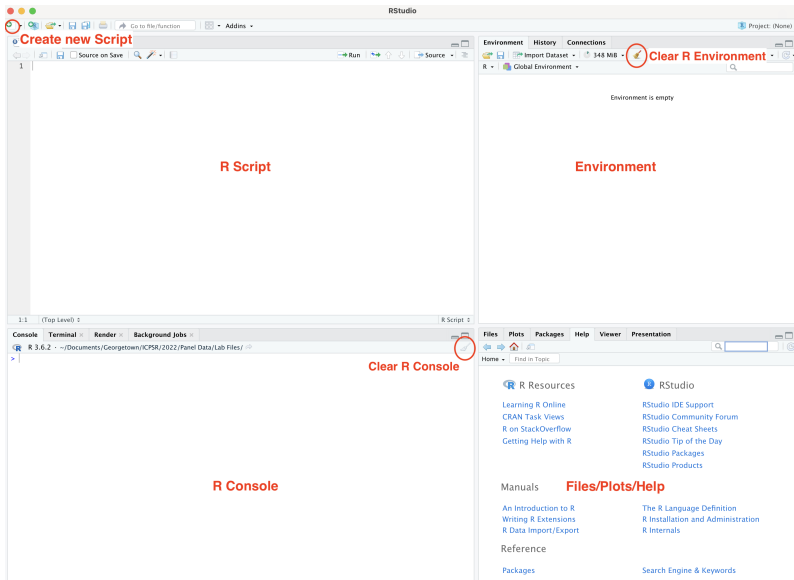
- Specify research question
- Develop research design
- Collect data
- Preprocess data
  - The data you have can contain errors, be incomplete, may need additional variable construction, etc.
  - Clean the dataset so that it can be used for analysis
  - Summarizing/Visualizing data
  - Statistical analysis
  - Summarizing/Visualizing results



# Basic Steps of Data Analysis

- Specify research question
- Develop research design
- Collect data
- Preprocess data
  - The data you have can contain errors, be incomplete, may need additional variable construction, etc.
  - Clean the dataset so that it can be used for analysis
  - Summarizing/Visualizing data
  - Statistical analysis
  - Summarizing/Visualizing results

# R Studio Interface



# R Studio Interface

- R Script
  - Write your code here

# R Studio Interface

- R Script
  - Write your code here
  - Each line is a separate command

# R Studio Interface

- R Script
  - Write your code here
  - Each line is a separate command
  - "Run" a line by highlighting it and clicking "Run" or pressing Cmd+Return (Mac) or Ctrl+Return (Windows)

# R Studio Interface

- R Script
  - Write your code here
  - Each line is a separate command
  - "Run" a line by highlighting it and clicking "Run" or pressing Cmd+Return (Mac) or Ctrl+Return (Windows)
  - Code should be written so humans can easily understand what's happening. Annotate/Comment your code by prefacing a line with #

# R Studio Interface

- R Script
  - Write your code here
  - Each line is a separate command
  - "Run" a line by highlighting it and clicking "Run" or pressing Cmd+Return (Mac) or Ctrl+Return (Windows)
  - Code should be written so humans can easily understand what's happening. Annotate/Comment your code by prefacing a line with #
  - The script can be saved as a .R file

# R Studio Interface

- R Script
  - Write your code here
  - Each line is a separate command
  - "Run" a line by highlighting it and clicking "Run" or pressing Cmd+Return (Mac) or Ctrl+Return (Windows)
  - Code should be written so humans can easily understand what's happening. Annotate/Comment your code by prefacing a line with #
  - The script can be saved as a .R file
- Console
  - Where results are displayed



# R Studio Interface

- R Script
  - Write your code here
  - Each line is a separate command
  - "Run" a line by highlighting it and clicking "Run" or pressing Cmd+Return (Mac) or Ctrl+Return (Windows)
  - Code should be written so humans can easily understand what's happening. Annotate/Comment your code by prefacing a line with `#`
  - The script can be saved as a .R file
- Console
  - Where results are displayed
  - You can type code directly into the console **but this is bad practice since the Console input cannot be saved**

# R Studio Interface

- R Script
  - Write your code here
  - Each line is a separate command
  - "Run" a line by highlighting it and clicking "Run" or pressing Cmd+Return (Mac) or Ctrl+Return (Windows)
  - Code should be written so humans can easily understand what's happening. Annotate/Comment your code by prefacing a line with `#`
  - The script can be saved as a .R file
- Console
  - Where results are displayed
  - You can type code directly into the console **but this is bad practice since the Console input cannot be saved**
- Environment
  - Where objects are stored

# R Studio Interface

- R Script
  - Write your code here
  - Each line is a separate command
  - "Run" a line by highlighting it and clicking "Run" or pressing Cmd+Return (Mac) or Ctrl+Return (Windows)
  - Code should be written so humans can easily understand what's happening. Annotate/Comment your code by prefacing a line with `#`
  - The script can be saved as a .R file
- Console
  - Where results are displayed
  - You can type code directly into the console **but this is bad practice since the Console input cannot be saved**
- Environment
  - Where objects are stored

# Object Class/Type/Mode

- character: non-numeric (text)

# Object Class/Type/Mode

- character: non-numeric (text)
- numeric: numbers

# Object Class/Type/Mode

- character: non-numeric (text)
- numeric: numbers
  - Includes "integer" and "double" — concerns level of precision; R works this out behind the scenes

# Object Class/Type/Mode

- character: non-numeric (text)
- numeric: numbers
  - Includes "integer" and "double" — concerns level of precision; R works this out behind the scenes
- logical: TRUE and FALSE

# Object Class/Type/Mode

- character: non-numeric (text)
- numeric: numbers
  - Includes "integer" and "double" — concerns level of precision; R works this out behind the scenes
- logical: TRUE and FALSE
- factor: a numeric which R has categorized into "levels"



# Object Class/Type/Mode

- character: non-numeric (text)
- numeric: numbers
  - Includes "integer" and "double" — concerns level of precision; R works this out behind the scenes
- logical: TRUE and FALSE
- factor: a numeric which R has categorized into "levels"

# Data Structures

- Vectors
  - Data analysis is not possible without data structures for data. R has several important data structures that shape how information is stored and processed.

# Data Structures

- Vectors
  - Data analysis is not possible without data structures for data. R has several important data structures that shape how information is stored and processed.
  - *Vectors* hold a sequence of values of the same data type, like numbers, characters, or logical values. They are one-dimensional arrays or lists of values that contain one and only one type of data. Vector types in R include: logical, integer, double, character, complex, raw.

# Data Structures

- Vectors
  - Data analysis is not possible without data structures for data. R has several important data structures that shape how information is stored and processed.
  - *Vectors* hold a sequence of values of the same data type, like numbers, characters, or logical values. They are one-dimensional arrays or lists of values that contain one and only one type of data. Vector types in R include: logical, integer, double, character, complex, raw.
- Matrices
  - *Matrices* are two-dimensional data structures in R that consists of rows and columns of values. Every element is of the same type. It's essentially a collection of vectors of the same length, organized in rows and columns.

# Data Structures

- Vectors
  - Data analysis is not possible without data structures for data. R has several important data structures that shape how information is stored and processed.
  - *Vectors* hold a sequence of values of the same data type, like numbers, characters, or logical values. They are one-dimensional arrays or lists of values that contain one and only one type of data. Vector types in R include: logical, integer, double, character, complex, raw.
- Matrices
  - *Matrices* are two-dimensional data structures in R that consists of rows and columns of values. Every element is of the same type. It's essentially a collection of vectors of the same length, organized in rows and columns.
  - Most data contains at least numeric information and character information. That is why we don't usually use matrices.

# Data Structures

- Data Frames
  - *Data Frames* and *Tibbles* are what you will usually use. They are a tabular data structure in R, similar to a spreadsheet or a database table. They can store different types of data (e.g., numeric, character) in columns. Each column is a vector, and all columns must have the same length.

# Data Structures

- Data Frames
  - *Data Frames* and *Tibbles* are what you will usually use. They are a tabular data structure in R, similar to a spreadsheet or a database table. They can store different types of data (e.g., numeric, character) in columns. Each column is a vector, and all columns must have the same length.
  - Most times, each column will be of one type while a given row will contain many different types. We usually refer to columns as *variables* and rows as *observations*.

# Data Structures

- Data Frames
  - *Data Frames* and *Tibbles* are what you will usually use. They are a tabular data structure in R, similar to a spreadsheet or a database table. They can store different types of data (e.g., numeric, character) in columns. Each column is a vector, and all columns must have the same length.
  - Most times, each column will be of one type while a given row will contain many different types. We usually refer to columns as *variables* and rows as *observations*.
- Example
  - Inspect R's built-in dataset, *mtcars*



# Objects and Assignment

- In R, to work with values/data, we *assign* them to *objects*

# Objects and Assignment

- In R, to work with values/data, we *assign* them to *objects*
  - Object: named "box" or "container" to store values/data. Think of it as a box where you can store numbers, text, or other types of information. You can give the box a name, and then you can use that name to access and manipulate the data inside.

# Objects and Assignment

- In R, to work with values/data, we *assign* them to *objects*
  - Object: named "box" or "container" to store values/data. Think of it as a box where you can store numbers, text, or other types of information. You can give the box a name, and then you can use that name to access and manipulate the data inside.
  - Using the assignment operator: `<-`

# Objects and Assignment

- In R, to work with values/data, we *assign* them to *objects*
  - Object: named "box" or "container" to store values/data. Think of it as a box where you can store numbers, text, or other types of information. You can give the box a name, and then you can use that name to access and manipulate the data inside.
  - Using the assignment operator: `<-`
    - "Less than" (Shift+comma) Hyphen

# Objects and Assignment

- In R, to work with values/data, we *assign* them to *objects*
  - Object: named "box" or "container" to store values/data. Think of it as a box where you can store numbers, text, or other types of information. You can give the box a name, and then you can use that name to access and manipulate the data inside.
  - Using the assignment operator: <-
    - "Less than" (Shift+comma) Hyphen
  - object <- values/data

# Objects and Assignment

- In R, to work with values/data, we *assign* them to *objects*
  - Object: named "box" or "container" to store values/data. Think of it as a box where you can store numbers, text, or other types of information. You can give the box a name, and then you can use that name to access and manipulate the data inside.
  - Using the assignment operator: `<-`
    - "Less than" (Shift+comma) Hyphen
  - `object <- values/data`
- [Exercise](#)

# Objects and Assignment

- In R, to work with values/data, we *assign* them to *objects*
  - Object: named "box" or "container" to store values/data. Think of it as a box where you can store numbers, text, or other types of information. You can give the box a name, and then you can use that name to access and manipulate the data inside.
  - Using the assignment operator: `<-`
    - "Less than" (Shift+comma) Hyphen
  - `object <- values/data`
- **Exercise**
  - Write two arithmetic operations in R and assign them to unique names. Then perform arithmetic operations using the named results.

# Objects and Assignment

- In R, to work with values/data, we *assign* them to *objects*
  - Object: named "box" or "container" to store values/data. Think of it as a box where you can store numbers, text, or other types of information. You can give the box a name, and then you can use that name to access and manipulate the data inside.
  - Using the assignment operator: `<-`
    - "Less than" (Shift+comma) Hyphen
  - `object <- values/data`
- **Exercise**
  - Write two arithmetic operations in R and assign them to unique names. Then perform arithmetic operations using the named results.

```
a <- 5 + 5 + 5
```

```
b <- 6 - 6 - 6
```

```
a + b
```



# Objects, Functions, and Arguments

- In R, we use *commands* (or *functions*) to perform tasks on these *objects*

# Objects, Functions, and Arguments

- In R, we use *commands* (or *functions*) to perform tasks on these *objects*
  - Data analysis requires a lot more than just basic arithmetic. R contains many *functions* that can perform mathematical operations, process data, run analyses, and more.

# Objects, Functions, and Arguments

- In R, we use *commands* (or *functions*) to perform tasks on these *objects*
  - Data analysis requires a lot more than just basic arithmetic. R contains many *functions* that can perform mathematical operations, process data, run analyses, and more.
  - A function is like a recipe that tells R to perform a specific task. You give the function some input (called arguments), and it gives you an output based on the task it's designed to do.

# Objects, Functions, and Arguments

- In R, we use *commands* (or *functions*) to perform tasks on these *objects*
  - Data analysis requires a lot more than just basic arithmetic. R contains many *functions* that can perform mathematical operations, process data, run analyses, and more.
  - A function is like a recipe that tells R to perform a specific task. You give the function some input (called arguments), and it gives you an output based on the task it's designed to do.
  - Type your function, then your object and other arguments in parentheses. **For Example:** the command `c(x, y, z)` will combine values `x`, `y`, and `z` into a vector.

# Objects, Functions, and Arguments

- In R, we use *commands* (or *functions*) to perform tasks on these *objects*
  - Data analysis requires a lot more than just basic arithmetic. R contains many *functions* that can perform mathematical operations, process data, run analyses, and more.
  - A function is like a recipe that tells R to perform a specific task. You give the function some input (called arguments), and it gives you an output based on the task it's designed to do.
  - Type your function, then your object and other arguments in parentheses. **For Example:** the command `c(x, y, z)` will combine values `x`, `y`, and `z` into a vector.
  - **Exercise:** Combine the 2 objects you created earlier into a vector called "values" using `c()`. Calculate the average of the values in the vector using the `mean()` function.

# Objects, Functions, and Arguments

- In R, we use *commands* (or *functions*) to perform tasks on these *objects*
  - Data analysis requires a lot more than just basic arithmetic. R contains many *functions* that can perform mathematical operations, process data, run analyses, and more.
  - A function is like a recipe that tells R to perform a specific task. You give the function some input (called arguments), and it gives you an output based on the task it's designed to do.
  - Type your function, then your object and other arguments in parentheses. **For Example:** the command `c(x, y, z)` will combine values `x`, `y`, and `z` into a vector.
  - **Exercise:** Combine the 2 objects you created earlier into a vector called "values" using `c()`. Calculate the average of the values in the vector using the `mean()` function.

```
values <- c(a, b)
mean(values)
```

# Objects, Functions, and Arguments

- R functions typically contain many *arguments*. Your object is usually one argument, other arguments can be thought of as "options".  
**Example:** `mean()` has `x`, `trim`, and `na.rm`.

# Objects, Functions, and Arguments

- R functions typically contain many *arguments*. Your object is usually one argument, other arguments can be thought of as "options".  
**Example:** `mean()` has `x`, `trim`, and `na.rm`.
- Many arguments have default values and don't need to be included in function calls.



# Objects, Functions, and Arguments

- R functions typically contain many *arguments*. Your object is usually one argument, other arguments can be thought of as "options".  
**Example:** `mean()` has `x`, `trim`, and `na.rm`.
- Many arguments have default values and don't need to be included in function calls.
- Functions are recognizable because they end in `()`. **Example:** `command(object, other arguments)`

# Objects, Functions, and Arguments

- R functions typically contain many *arguments*. Your object is usually one argument, other arguments can be thought of as "options".  
**Example:** `mean()` has `x`, `trim`, and `na.rm`.
- Many arguments have default values and don't need to be included in function calls.
- Functions are recognizable because they end in `()`. **Example:** `command(object, other arguments)`
- Documentation for functions can be easily accessed by prefacing the function name with `?` and dropping the `()`. The documentation typically includes a description, a list of the arguments, references, a list of related functions, and examples. The examples are incredibly useful. **Example:** `?mean`

# Packages

- Opening RStudio automatically loads base R, a fundamental collection of code and functions.

# Packages

- Opening RStudio automatically loads base R, a fundamental collection of code and functions.
- R is an extensible programming language. Users contribute millions of lines of code that can be used by other users without condition or compensation. The main mode of contribution are R packages.

# Packages

- Opening RStudio automatically loads base R, a fundamental collection of code and functions.
- R is an extensible programming language. Users contribute millions of lines of code that can be used by other users without condition or compensation. The main mode of contribution are R packages.
- A package is a collection of functions, data, and documentations which is publicly shared to enhance the functionality of R.

# Packages

- Opening RStudio automatically loads base R, a fundamental collection of code and functions.
- R is an extensible programming language. Users contribute millions of lines of code that can be used by other users without condition or compensation. The main mode of contribution are R packages.
- A package is a collection of functions, data, and documentations which is publicly shared to enhance the functionality of R.

# Packages

- Installing packages

# Packages

- Installing packages
  - Packages can be installed using `install.packages("")`. Include the name of the desired package in quotes as the only argument to the function.



# Packages

- Installing packages
  - Packages can be installed using `install.packages("")`. Include the name of the desired package in quotes as the only argument to the function.
  - **Example:** Install the package, *tidyverse*

# Packages

- Installing packages
  - Packages can be installed using `install.packages("")`. Include the name of the desired package in quotes as the only argument to the function.
  - **Example:** Install the package, *tidyverse*
  - It is customary to never include `install.packages()` in a .R script.

# Packages

- Installing packages
  - Packages can be installed using `install.packages("")`. Include the name of the desired package in quotes as the only argument to the function.
  - **Example:** Install the package, *tidyverse*
  - It is customary to never include `install.packages()` in a .R script.
- Using packages
  - After installation, packages need to be loaded once per R session using the `library()` function. Include the name of the desired package without quotes as the only argument to the function.

# Packages

- Installing packages
  - Packages can be installed using `install.packages("")`. Include the name of the desired package in quotes as the only argument to the function.
  - **Example:** Install the package, *tidyverse*
  - It is customary to never include `install.packages()` in a .R script.
- Using packages
  - After installation, packages need to be loaded once per R session using the `library()` function. Include the name of the desired package without quotes as the only argument to the function.
  - **Example:** Load the package, *tidyverse*

# Packages

- Installing packages
  - Packages can be installed using `install.packages("")`. Include the name of the desired package in quotes as the only argument to the function.
  - **Example:** Install the package, *tidyverse*
  - It is customary to never include `install.packages()` in a .R script.
- Using packages
  - After installation, packages need to be loaded once per R session using the `library()` function. Include the name of the desired package without quotes as the only argument to the function.
  - **Example:** Load the package, *tidyverse*
  - It is a good idea to include `library()` statements at the top of scripts for each package used in the script. This way it is obvious at the top of the script which packages are necessary.